

Linux Assembly

Armine Hayrapetyan

Outline

- Introduction to Files
- Open Files
- Writing Files
- Reading Files

File Permissions

Files modes specify the *permissions* for files. Essentially, they specify who is allowed to read, write, and/or execute the file.

Modes are stored as a three-digit octal value (base-8).

Octal Value Permissions

Value	Read	Write	Execute
0			
1			X
2		X	
3		X	X
4	X		
5	X		X
6	X	X	
7	X	X	X

File Permissions

File permissions in Linux are set with four octal values. The least 3 significant octal values are for the file owner's permissions, the group's permissions, and the "other's" permissions (those outside the group and not the file owner).

The most significant octal value is reserved for special permissions.

Usually we do not need to use these.

	Special	Owner	Group	Other
1	sticky bit	execute	execute	execute
2	setgid	write	write	write
4	setuid	read	read	read

System Calls

We have previously used `sys_read` and `sys_write` to read text from and write text to the standard input and output respectively.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
<code>sys_read</code>	0	<code>#filedescriptor</code>	<code>\$buffer</code>	<code>#count</code>			
<code>sys_write</code>	1	<code>#filedescriptor</code>	<code>\$buffer</code>	<code>#count</code>			
<code>sys_open</code>	2	<code>\$filename</code>	<code>#flags</code>	<code>#mode</code>			
<code>sys_close</code>	3	<code>#filedescriptor</code>					
...
<code>pwritev2</code>	328

System Calls

`sys_close` is used when a file is no longer in use.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
<code>sys_read</code>	0	<code>#filedescriptor</code>	<code>\$buffer</code>	<code>#count</code>			
<code>sys_write</code>	1	<code>#filedescriptor</code>	<code>\$buffer</code>	<code>#count</code>			
<code>sys_open</code>	2	<code>\$filename</code>	<code>#flags</code>	<code>#mode</code>			
<code>sys_close</code>	3	<code>#filedescriptor</code>					
...
<code>pwritev2</code>	328

sys_open

The first argument sys_open takes is a pointer to the filename (zero terminated).

The second argument are the flags.

The third argument is the file mode, being the 4-digit octal number that we learned from earlier.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
sys_open	2	\$filename	#flags	#mode			

sys_open

Here is the code to open a file with the “create” and “write” flag.

The “create” flag creates the file if it does not exist.

```
mov rax, SYS_OPEN
mov rdi, filename
mov rsi, O_CREAT+O_WRONLY
mov rdx, 0644o
syscall
```


sys_open

This is the ID of the system call, specifically the ID for sys_open.

This is the pointer to the zero-terminated string for the file name to open.

These are the “create” (64) and “write” (1) flags.

These are the file permissions we learned earlier.

```
mov rax, SYS_OPEN
mov rdi, filename
mov rsi, 0_CREAT+0_WRONLY
mov rdx, 0644o
syscall
```

The “o” tells NASM this is an octal value.

sys_open

This system call returns the file descriptor of the file opened within the rax register.
If there is an error, that error is returned in the rax register.

sys_write

This system call can be used to write text to a file.

It is used exactly like how we used it in the “Hello, World!” video to display text on the screen. The only difference is that the first argument is changed to the file descriptor returned from the `sys_open` system call.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
sys_write	1	#filedescriptor	\$buffer	#count			

sys_write

Here is code to write text to a file opened.

```
mov rdi, rax  
mov rax, SYS_WRITE  
mov rsi, text  
mov rdx, 17  
syscall
```

sys_write

The file descriptor comes from the rax register assuming sys_open was successful.

This is the ID of the system call, specifically the ID for sys_write.

A pointer to the text which will be written to the file.

The number of bytes to write to the file, in this case 17 bytes.

```
mov rdi, rax
mov rax, SYS_WRITE
mov rsi, text
mov rdx, 17
syscall
```

SYS_WRITE equ 1

sys_close

sys_close only takes the file descriptor as its only argument.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
sys_close	3	#filedescriptor					

sys_close

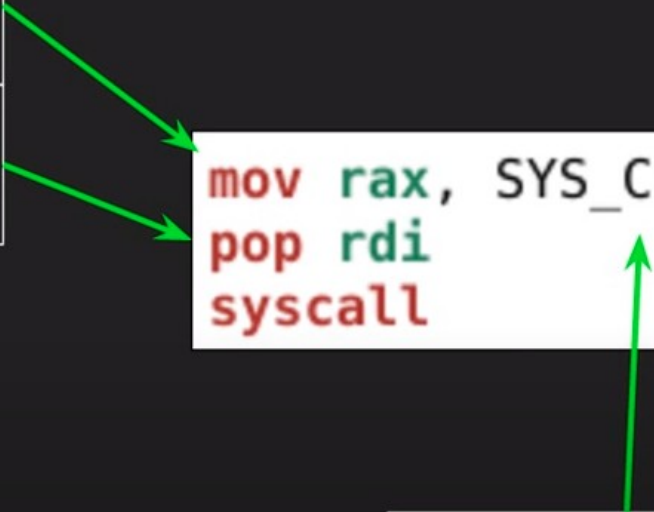
Here is code to close an opened file.

```
mov rax, SYS_CLOSE  
pop rdi  
syscall
```

sys_close

This is the ID of the system call, specifically the ID for sys_close.

This is the file descriptor of the file to close, it assumes it is on the top of the stack.



```
mov rax, SYS_CLOSE  
pop rdi  
syscall
```

SYS_CLOSE equ 3

Code to Write to a File

Find this code here:

<http://pastebin.com/b9TUxfzg>

Open the file

Write to the file

Close the file

```
%include "linux64.inc"

section .data
    filename db "myfile.txt",0
    text db "Here's some text."

section .text
    global _start
_start:
    mov rax, SYS_OPEN
    mov rdi, filename
    mov rsi, 0_CREAT+0_WRONLY
    mov rdx, 0644o
    syscall

    push rax
    mov rdi, rax
    mov rax, SYS_WRITE
    mov rsi, text
    mov rdx, 17
    syscall

    mov rax, SYS_CLOSE
    pop rdi
    syscall

    exit
```

sys_open

The first argument `sys_open` takes is a pointer to the filename (zero terminated).

The second argument are the flags.

The third argument is the file mode, being the 4-digit octal number that we learned from earlier.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
sys_open	2	<code>\$filename</code>	<code>#flags</code>	<code>#mode</code>			

sys_open

Flag Name	Value	$\log_2(\text{value})$
O_RDONLY	0	null
O_WRONLY	1	0
O_RDWR	2	1
O_CREAT	64	6
O_APPEND	1024	10
O_DIRECTORY	65536	16
O_PATH	2097152	21
O_TMPFILE	4194304	22

sys_open

Here is the code to open a file with the “read only” flag.

```
mov rax, SYS_OPEN  
mov rdi, filename  
mov rsi, O_RDONLY  
mov rdx, 0644o  
syscall
```

sys_open

This is the ID of the system call, specifically the ID for sys_open.

This is the pointer to the zero-terminated string for the file name to open.

This is the “read” flag (0).

This is the file permission, but it does not matter if we are only reading the file.

```
mov rax, SYS_OPEN  
mov rdi, filename  
mov rsi, 0_RDONLY  
mov rdx, 0644o  
syscall
```

The “o” tells NASM this is an octal value.

sys_open

This system call returns the file descriptor of the file opened within the rax register. If there is an error, that error is returned in the rax register.

sys_read

This system call can be used to read text from a file.

It is used exactly like how we used it in the tutorial on getting user input. The only difference is that the first argument is changed to the file descriptor returned from the `sys_open` system call.

syscall	ID	ARG1	ARG2	ARG3	ARG4	ARG5	ARG6
sys_read	0	#filedescriptor	\$buffer	#count			

sys_read

Here is code to read from an opened file.

```
mov rdi, rax  
mov rax, SYS_READ  
mov rsi, text  
mov rdx, 17  
syscall
```


sys_read

The file descriptor comes from the `rax` register assuming `sys_open` was successful.

This is the ID of the system call, specifically the ID for `sys_read`.

A pointer to where the read text will be stored.

The number of bytes to read from the file, in this case 17 bytes.

```
mov rdi, rax
mov rax, SYS_READ
mov rsi, text
mov rdx, 17
syscall
```

`SYS_READ equ 0`

sys_close

This is the ID of the system call, specifically the ID for sys_close.

This is the file descriptor of the file to close, it assumes it is on the top of the stack.

```
mov rax, SYS_CLOSE  
pop rdi  
syscall
```

SYS_CLOSE equ 3

Code to Write to a File

Find this code here:

<http://pastebin.com/xcFtXk3t>

Open the file

Read from the file

Close the file

```
%include "linux64.inc"

section .data
    filename db "myfile.txt",0

section .bss
    text resb 18

section .text
    global _start
_start:
    mov rax, SYS_OPEN
    mov rdi, filename
    mov rsi, 0_RDONLY
    mov rdx, 0
    syscall

    push rax
    mov rdi, rax
    mov rax, SYS_READ
    mov rsi, text
    mov rdx, 17
    syscall

    mov rax, SYS_CLOSE
    pop rdi
    syscall

    print text
    exit
```

References

- <https://www.youtube.com/watch?v=AwmhZUATGYM&list=PLetF-YjXm-sCH6FrTz4AQhfH6INDQvQSn&index=10>
- https://www.youtube.com/watch?v=vXsUIX_Ozgc&list=PLetF-YjXm-sCH6FrTz4AQhfH6INDQvQSn&index=11
- https://www.youtube.com/watch?v=vXsUIX_Ozgc&list=PLetF-YjXm-sCH6FrTz4AQhfH6INDQvQSn&index=12